

# Lab 1

Conie O'Malley

2025-01-21

```
# libraries
library(reticulate)
```

Warning: package 'reticulate' was built under R version 4.3.3

```
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
v dplyr      1.1.4      v readr      2.1.5
v forcats    1.0.0      v stringr    1.5.1
v ggplot2     3.5.1      v tibble     3.2.1
v lubridate  1.9.3      v tidyr      1.3.1
v purrr       1.0.2
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()     masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become
```

```
library(ggplot2)
#use_condaenv("datascience", required = FALSE) # set my environment
```

## Part 1: Hello World and Beyond

### Issuing Interactive Commands & Adding Comments

```
print("Hello World!") # first line r of code
```

```
[1] "Hello World!"
```

```
print("Hello Stats!") # second line r of code
```

```
[1] "Hello Stats!"
```

```
print("Hello World!") # frist line of python code
```

```
Hello World!
```

```
print("Hello Stats!") # second line of python code
```

```
Hello Stats!
```

## Doing Simple Math Calculations

```
1 + 2 + 3 + 4 + 5
```

```
[1] 15
```

```
sum(1:5) # alternate way to write the code
```

```
[1] 15
```

## Creating and Using Vectors and Operations

```
c(1, 2, 3, 4, 5) # concatenate
```

```
[1] 1 2 3 4 5
```

```
1:5 # sequence operator
```

```
[1] 1 2 3 4 5
```

```
sum(1:5) # addition using the sequence operator
```

```
[1] 15
```

## Storing and Calculating Values

```
x <- 1:5 # vector assignment  
y <- 10 # vector assignment  
x+y # vector addition
```

```
[1] 11 12 13 14 15
```

```
z <- x+y # vector assignment  
z
```

```
[1] 11 12 13 14 15
```

```
h <- "Hello" # vector assignment  
h
```

```
[1] "Hello"
```

```
hw <- c("Hello", "World!") # vector concatenation  
print(hw) # print view
```

```
[1] "Hello" "World!"
```

```
paste(hw) # paste view
```

```
[1] "Hello" "World!"
```

## Navigating the RStudio Workspace

```
ls() # view created objects
```

```
[1] "h" "hw" "x" "y" "z"
```

```
rm("z") # remove "z" object  
ls() # confirm removal
```

```
[1] "h" "hw" "x" "y"
```

## More Practice Vectorizing & Vectors of Unequal Length

```
baskets.of.granny <- c(12, 4, 4, 6, 9, 3)  
sum(baskets.of.granny)
```

```
[1] 38
```

```
firstnames <- c("John", "Jacqueline", "Robert") # create vector list  
lastname <- "Kennedy" # create vector  
paste(firstnames, lastname) # paste vectors
```

```
[1] "John Kennedy" "Jacqueline Kennedy" "Robert Kennedy"
```


```
lastnames <- c("Kennedy", "Kennedy-Onnasis") # create vector list  
paste(firstnames, lastnames) # paste vectors
```

```
[1] "John Kennedy" "Jacqueline Kennedy-Onnasis"  
[3] "Robert Kennedy"
```

## Part 2: Statistical Analysis with R

### Scatter Plot

```
ggplot(mpg, aes(displ, hwy)) + # create a scatter plot
  geom_point() +
  labs(title = "Displacement vs. Highway MPG Scatterplot",
       x = "Displacement (litres)",
       y = "Highway MPG") + # apply labels
  theme(plot.title = element_text(hjust = 0.5))
```




## Analysis

There is an inverse relationship between the volume of engine displacement and highway miles per gallon. The relationship appears to be slightly curvilinear, but cannot be confirmed without a residuals plot.

## Box Plot

```
ggplot(mpg, aes(class, hwy)) + # create a box plot
  geom_boxplot() +
  labs(title = "Class vs. Highway MPG Box Plot",
       x = "Class",
       y = "Highway MPG") + # apply labels
  theme(plot.title = element_text(hjust = 0.5))
```



### Analysis

There seems to be some commonality among the types *2seater*, *compact*, *midsize*, and *subcompact*, because they all have overlapping box plot ranges. *pickup* and *suv* also fall into their own common area because of overlapping ranges, while the *minivan* class has no overlaps. There are a considerable amount of outliers in the *suv* class - meaning that, at first glance, this is the most varied class of vehicles, where as *2seater* has no outliers and a very tight range - meaning that, at first glance, this is likely the most homogeneous class of vehicles.

### Computing Basic Statistics

```
mean(economics$unemploy) # calculate mean
```

```
[1] 7771.31
```

```
var(economics$unemploy) # calculate variance
```

```
[1] 6979948
```

```
sd(economics$unemploy) # calculate standard deviation
```

```
[1] 2641.959
```

```
min(economics$unemploy) # calculate min
```

```
[1] 2685
```

```
max(economics$unemploy) # calculate max
```

```
[1] 15352
```

```
median(economics$unemploy) # calculate median
```

```
[1] 7494
```

```
cor(economics$pce, economics$psavert) # calculate correlation
```

```
[1] -0.7928546
```

```
pce_psavert_cor <- round(cor(economics$pce, economics$psavert),4)  
# assign correlation to vector variable
```

## Analysis

There is a strong, negative correlation between *pce* and *psavert* (-0.7929).

## Conducting a t-test

```
data(tips, package = 'reshape2') # attach data  
t.test(tips$tip, alternative = 'two.sided', mu=2.50) # conduct two tail t-test
```

### One Sample t-test

```
data: tips$tip
t = 5.6253, df = 243, p-value = 5.08e-08
alternative hypothesis: true mean is not equal to 2.5
95 percent confidence interval:
 2.823799 3.172758
sample estimates:
mean of x
 2.998279
```

### Analysis

$t$  is our t-value, which is the standardized test statistic for this data set. Our t-value should be greater than our t statistic we are testing against, so we can reject  $H_0 : \mu = 2.50$ .  $df$  are the degrees of freedom in this data set.  $p$ -value is the probability that we will get an sample mean under the  $H_0$ . Our p-value ( $5.08e-08$ )  $< 0.05$ , meaning we can reject  $H_0$ . *confidence interval* is 95% - meaning that if we sampled the data randomly, our sample mean would be within the range 95% of the time. *sample mean of  $x$*  is the mean of our current sample from the data. We can reject  $H_0 : \mu = 2.50$  because our p-value  $< 0.05$ .

### Building a Linear Regression Model

```
head(mpg) # view first 6 rows of data
```

```
# A tibble: 6 x 11
  manufacturer model displ  year   cyl trans      drv   cty   hwy fl   class
  <chr>         <chr> <dbl> <int> <int> <chr>   <chr> <int> <int> <chr> <chr>
1 audi         a4      1.8  1999     4 auto(l5) f      18    29 p   compa~
2 audi         a4      1.8  1999     4 manual(m5) f      21    29 p   compa~
3 audi         a4      2    2008     4 manual(m6) f      20    31 p   compa~
4 audi         a4      2    2008     4 auto(av) f      21    30 p   compa~
5 audi         a4      2.8  1999     6 auto(l5) f      16    26 p   compa~
6 audi         a4      2.8  1999     6 manual(m5) f      18    26 p   compa~
```

```
tail(mpg) # view last 6 rows of data
```



```
# A tibble: 6 x 11
  manufacturer model displ year   cyl trans      drv   cty   hwy fl      class
  <chr>          <chr> <dbl> <int> <int> <chr>    <chr> <int> <int> <chr> <chr>
1 volkswagen  passat   1.8  1999     4 auto(l5)  f       18    29 p    mids~
2 volkswagen  passat   2    2008     4 auto(s6)  f       19    28 p    mids~
3 volkswagen  passat   2    2008     4 manual(m6) f       21    29 p    mids~
4 volkswagen  passat   2.8  1999     6 auto(l5)  f       16    26 p    mids~
5 volkswagen  passat   2.8  1999     6 manual(m5) f       18    26 p    mids~
6 volkswagen  passat   3.6  2008     6 auto(s6)  f       17    26 p    mids~
```

```
lm(hwy ~ displ, mpg) # build a basic linear regression model
```

Call:


```
lm(formula = hwy ~ displ, data = mpg)
```

Coefficients:

```
(Intercept)      displ
    35.698         -3.531
```

```
ggplot(mpg, aes(displ, hwy)) + # build a linear regression model scatter plot
  geom_point() +
  labs(title = "Displacement vs. Highway MPG Regression Model",
       x = "Engine Size",
       y = "Highway Fuel Efficiency") +
  geom_smooth(method = "lm", color = "red") +
  theme(plot.title = element_text(hjust = 0.5))
```

```
`geom_smooth()` using formula = 'y ~ x'
```



```
fuelILM <- lm(displ ~ hwy, mpg) # assign new model to vector variable
fuelILM
```

Call:

```
lm(formula = displ ~ hwy, data = mpg)
```

Coefficients:

(Intercept)	hwy
7.3676	-0.1662

```
summary(fuelILM) # review summary statistics
```

Call:

```
lm(formula = displ ~ hwy, data = mpg)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.4126	-0.5710	-0.1105	0.4571	3.6212

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	7.367570	0.221422	33.27	<2e-16 ***
hwy	-0.166201	0.009157	-18.15	<2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8323 on 232 degrees of freedom

Multiple R-squared: 0.5868, Adjusted R-squared: 0.585

F-statistic: 329.5 on 1 and 232 DF, p-value: < 2.2e-16

## Part 3: Basic Importing and Wrangling of Data

### Inspecting and Cleaning the Data

```
housing <- read.table("http://www.jaredlander.com/data/housing.csv", sep = ",",
                      header = TRUE, stringsAsFactors = FALSE) # read data
names(housing) <- c("Neighborhood", "Class", "Units", "YearsBuilt", "SqFt",
                    "Income", "IncomePerSqFt", "Expense", "ExpensePerSqFt",
                    "NetIncome", "Value", "ValuePerSqFt", "Boro") # rename columns
head(housing)
```

	Neighborhood	Class	Units	YearsBuilt	SqFt	Income	IncomePer-SqFt
1	FINANCIAL	R9-CONDOMINIUM	42	1920	36500	1332615	36.51
2	FINANCIAL	R4-CONDOMINIUM	78	1985	126420	6633257	52.47
3	FINANCIAL	RR-CONDOMINIUM	500	NA	554174	17310000	31.24
4	FINANCIAL	R4-CONDOMINIUM	282	1930	249076	11776313	47.28
5	TRIBECA	R4-CONDOMINIUM	239	1985	219495	10004582	45.58
6	TRIBECA	R4-CONDOMINIUM	133	1986	139719	5127687	36.70

	Expense	ExpensePerSqFt	NetIncome	Value	ValuePerSqFt	Boro
1	342005	9.37	990610	7300000	200.00	Manhattan
2	1762295	13.94	4870962	30690000	242.76	Manhattan
3	3543000	6.39	13767000	90970000	164.15	Manhattan
4	2784670	11.18	8991643	67556006	271.23	Manhattan
5	2783197	12.68	7221385	54320996	247.48	Manhattan
6	1497788	10.72	3629899	26737996	191.37	Manhattan

### Building a Linear Regression Model

```
house1 <- lm(ValuePerSqFt ~ Units + SqFt + Boro,
             housing) # build a linear regression model
summary(house1) # view summary statistics
```

Call:

```
lm(formula = ValuePerSqFt ~ Units + SqFt + Boro, data = housing)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-164.418	-22.692	1.416	26.972	261.122

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	4.329e+01	5.330e+00	8.122	6.97e-16	***
Units	-1.881e-01	2.210e-02	-8.511	< 2e-16	***
SqFt	2.103e-04	2.087e-05	10.079	< 2e-16	***
BoroBrooklyn	3.456e+01	5.535e+00	6.244	4.95e-10	***
BoroManhattan	1.310e+02	5.385e+00	24.327	< 2e-16	***
BoroQueens	3.299e+01	5.663e+00	5.827	6.35e-09	***
BoroStaten Island	-3.630e+00	9.993e+00	-0.363	0.716	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 43.35 on 2619 degrees of freedom

Multiple R-squared: 0.6009, Adjusted R-squared: 0.6

F-statistic: 657.2 on 6 and 2619 DF, p-value: < 2.2e-16

## Part 4: Hello World, Data, Statistics and Beyond in Python

### Hello World in Python

```
print("Hello World!") # print function practice
```

Hello World!

```
print("Hello Stats!") # print function practice
```

Hello Stats!

```
print("Hello", "World!") # print function practice
```

Hello World!

## Doing Simple Math Calculations

```
1+2
```

3

```
1+2+3+4+5
```

15

## Installing and importing packages

```
from matplotlib import pyplot as plt # import libraries
from statsmodels.formula.api import ols
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn import datasets
from sklearn.datasets import fetch_california_housing
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import pandas as pd
import seaborn as sns
```

## Accessing a Built-In Dataset with Python

```
housing = fetch_california_housing() # assign data to matrix vector
X,y = housing.data, housing.target # assign subsets to variable vector
print("The size of the dataset is {}".format(X.shape))
```

The size of the dataset is (20640, 8)

```
print("The names of the data columns are {}", housing.feature_names)
```

The names of the data columns are {} ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Popula

```
print(housing.keys())
```

```
dict_keys(['data', 'target', 'frame', 'target_names', 'feature_names', 'DESCR'])
```

```
hypothesis = LinearRegression() # set hypothesis model type
hypothesis.fit(X,y) # fit the model
```

```
LinearRegression()
```

```
print(hypothesis.coef_) # print coefficients
```

```
[ 4.36693293e-01  9.43577803e-03 -1.07322041e-01  6.45065694e-01
 -3.97638942e-06 -3.78654265e-03 -4.21314378e-01 -4.34513755e-01]
```

## Accessing and Exploring Another Built-in Dataset in Python

```
iris = datasets.load_iris() # assign data set to matrix vector
iris.keys() # view keys
```

```
dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR', 'feature_names', 'filename',
```

```
iris['data'] # view data
```

```
array([[5.1, 3.5, 1.4, 0.2],
       [4.9, 3. , 1.4, 0.2],
       [4.7, 3.2, 1.3, 0.2],
       [4.6, 3.1, 1.5, 0.2],
       [5. , 3.6, 1.4, 0.2],
       [5.4, 3.9, 1.7, 0.4],
       [4.6, 3.4, 1.4, 0.3],
```

[5. , 3.4, 1.5, 0.2],  
[4.4, 2.9, 1.4, 0.2],  
[4.9, 3.1, 1.5, 0.1],  
[5.4, 3.7, 1.5, 0.2],  
[4.8, 3.4, 1.6, 0.2],  
[4.8, 3. , 1.4, 0.1],  
[4.3, 3. , 1.1, 0.1],  
[5.8, 4. , 1.2, 0.2],  
[5.7, 4.4, 1.5, 0.4],  
[5.4, 3.9, 1.3, 0.4],  
[5.1, 3.5, 1.4, 0.3],  
[5.7, 3.8, 1.7, 0.3],  
[5.1, 3.8, 1.5, 0.3],  
[5.4, 3.4, 1.7, 0.2],  
[5.1, 3.7, 1.5, 0.4],  
[4.6, 3.6, 1. , 0.2],  
[5.1, 3.3, 1.7, 0.5],  
[4.8, 3.4, 1.9, 0.2],  
[5. , 3. , 1.6, 0.2],  
[5. , 3.4, 1.6, 0.4],  
[5.2, 3.5, 1.5, 0.2],  
[5.2, 3.4, 1.4, 0.2],  
[4.7, 3.2, 1.6, 0.2],  
[4.8, 3.1, 1.6, 0.2],  
[5.4, 3.4, 1.5, 0.4],  
[5.2, 4.1, 1.5, 0.1],  
[5.5, 4.2, 1.4, 0.2],  
[4.9, 3.1, 1.5, 0.2],  
[5. , 3.2, 1.2, 0.2],  
[5.5, 3.5, 1.3, 0.2],  
[4.9, 3.6, 1.4, 0.1],  
[4.4, 3. , 1.3, 0.2],  
[5.1, 3.4, 1.5, 0.2],  
[5. , 3.5, 1.3, 0.3],  
[4.5, 2.3, 1.3, 0.3],  
[4.4, 3.2, 1.3, 0.2],  
[5. , 3.5, 1.6, 0.6],  
[5.1, 3.8, 1.9, 0.4],  
[4.8, 3. , 1.4, 0.3],  
[5.1, 3.8, 1.6, 0.2],  
[4.6, 3.2, 1.4, 0.2],  
[5.3, 3.7, 1.5, 0.2],  
[5. , 3.3, 1.4, 0.2],

[7. , 3.2, 4.7, 1.4],  
[6.4, 3.2, 4.5, 1.5],  
[6.9, 3.1, 4.9, 1.5],  
[5.5, 2.3, 4. , 1.3],  
[6.5, 2.8, 4.6, 1.5],  
[5.7, 2.8, 4.5, 1.3],  
[6.3, 3.3, 4.7, 1.6],  
[4.9, 2.4, 3.3, 1. ],  
[6.6, 2.9, 4.6, 1.3],  
[5.2, 2.7, 3.9, 1.4],  
[5. , 2. , 3.5, 1. ],  
[5.9, 3. , 4.2, 1.5],  
[6. , 2.2, 4. , 1. ],  
[6.1, 2.9, 4.7, 1.4],  
[5.6, 2.9, 3.6, 1.3],  
[6.7, 3.1, 4.4, 1.4],  
[5.6, 3. , 4.5, 1.5],  
[5.8, 2.7, 4.1, 1. ],  
[6.2, 2.2, 4.5, 1.5],  
[5.6, 2.5, 3.9, 1.1],  
[5.9, 3.2, 4.8, 1.8],  
[6.1, 2.8, 4. , 1.3],  
[6.3, 2.5, 4.9, 1.5],  
[6.1, 2.8, 4.7, 1.2],  
[6.4, 2.9, 4.3, 1.3],  
[6.6, 3. , 4.4, 1.4],  
[6.8, 2.8, 4.8, 1.4],  
[6.7, 3. , 5. , 1.7],  
[6. , 2.9, 4.5, 1.5],  
[5.7, 2.6, 3.5, 1. ],  
[5.5, 2.4, 3.8, 1.1],  
[5.5, 2.4, 3.7, 1. ],  
[5.8, 2.7, 3.9, 1.2],  
[6. , 2.7, 5.1, 1.6],  
[5.4, 3. , 4.5, 1.5],  
[6. , 3.4, 4.5, 1.6],  
[6.7, 3.1, 4.7, 1.5],  
[6.3, 2.3, 4.4, 1.3],  
[5.6, 3. , 4.1, 1.3],  
[5.5, 2.5, 4. , 1.3],  
[5.5, 2.6, 4.4, 1.2],  
[6.1, 3. , 4.6, 1.4],  
[5.8, 2.6, 4. , 1.2],



[5. , 2.3, 3.3, 1. ],  
 [5.6, 2.7, 4.2, 1.3],  
 [5.7, 3. , 4.2, 1.2],  
 [5.7, 2.9, 4.2, 1.3],  
 [6.2, 2.9, 4.3, 1.3],  
 [5.1, 2.5, 3. , 1.1],  
 [5.7, 2.8, 4.1, 1.3],  
 [6.3, 3.3, 6. , 2.5],  
 [5.8, 2.7, 5.1, 1.9],  
 [7.1, 3. , 5.9, 2.1],  
 [6.3, 2.9, 5.6, 1.8],  
 [6.5, 3. , 5.8, 2.2],  
 [7.6, 3. , 6.6, 2.1],  
 [4.9, 2.5, 4.5, 1.7],  
 [7.3, 2.9, 6.3, 1.8],  
 [6.7, 2.5, 5.8, 1.8],  
 [7.2, 3.6, 6.1, 2.5],  
 [6.5, 3.2, 5.1, 2. ],  
 [6.4, 2.7, 5.3, 1.9],  
 [6.8, 3. , 5.5, 2.1],  
 [5.7, 2.5, 5. , 2. ],  
 [5.8, 2.8, 5.1, 2.4],  
 [6.4, 3.2, 5.3, 2.3],  
 [6.5, 3. , 5.5, 1.8],  
 [7.7, 3.8, 6.7, 2.2],  
 [7.7, 2.6, 6.9, 2.3],  
 [6. , 2.2, 5. , 1.5],  
 [6.9, 3.2, 5.7, 2.3],  
 [5.6, 2.8, 4.9, 2. ],  
 [7.7, 2.8, 6.7, 2. ],  
 [6.3, 2.7, 4.9, 1.8],  
 [6.7, 3.3, 5.7, 2.1],  
 [7.2, 3.2, 6. , 1.8],  
 [6.2, 2.8, 4.8, 1.8],  
 [6.1, 3. , 4.9, 1.8],  
 [6.4, 2.8, 5.6, 2.1],  
 [7.2, 3. , 5.8, 1.6],  
 [7.4, 2.8, 6.1, 1.9],  
 [7.9, 3.8, 6.4, 2. ],  
 [6.4, 2.8, 5.6, 2.2],  
 [6.3, 2.8, 5.1, 1.5],  
 [6.1, 2.6, 5.6, 1.4],  
 [7.7, 3. , 6.1, 2.3],

```
[6.3, 3.4, 5.6, 2.4],
[6.4, 3.1, 5.5, 1.8],
[6. , 3. , 4.8, 1.8],
[6.9, 3.1, 5.4, 2.1],
[6.7, 3.1, 5.6, 2.4],
[6.9, 3.1, 5.1, 2.3],
[5.8, 2.7, 5.1, 1.9],
[6.8, 3.2, 5.9, 2.3],
[6.7, 3.3, 5.7, 2.5],
[6.7, 3. , 5.2, 2.3],
[6.3, 2.5, 5. , 1.9],
[6.5, 3. , 5.2, 2. ],
[6.2, 3.4, 5.4, 2.3],
[5.9, 3. , 5.1, 1.8]])
```

```
iris_df = pd.DataFrame(iris.data, columns = iris.feature_names) # create a df
iris_df['species'] = pd.Categorical.from_codes(iris.target, # add species column
                                              iris.target_names)
print(iris_df.head())
```

	sepal length (cm)	sepal width (cm)	...	petal width (cm)	species
0	5.1	3.5	...	0.2	setosa
1	4.9	3.0	...	0.2	setosa
2	4.7	3.2	...	0.2	setosa
3	4.6	3.1	...	0.2	setosa
4	5.0	3.6	...	0.2	setosa

[5 rows x 5 columns]

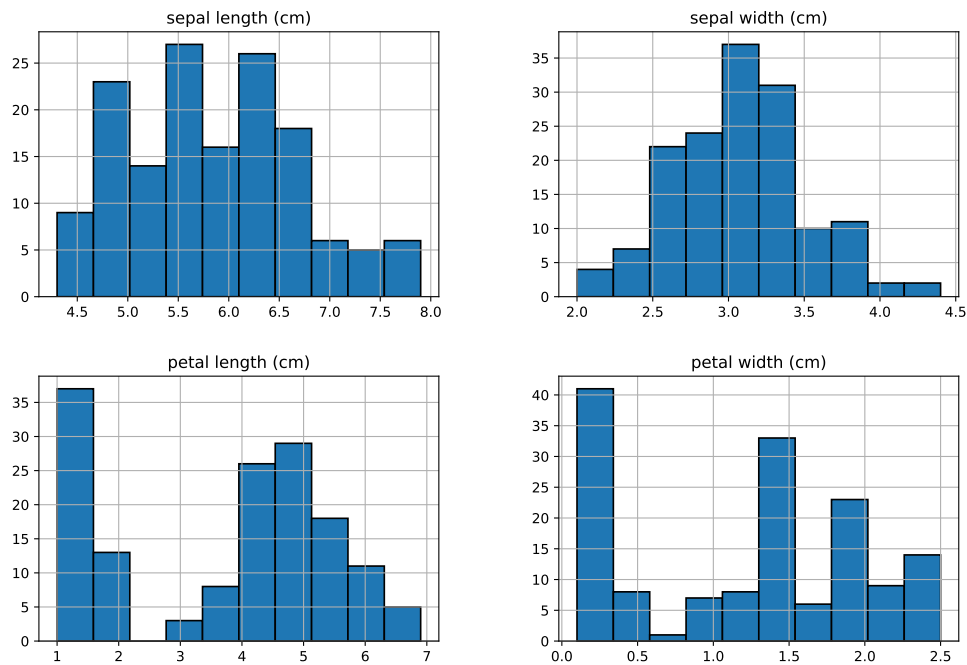
```
print(iris_df.describe())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333
std	0.828066	0.435866	1.765298	0.762238
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000


```
iris_df.hist(edgecolor = 'black', linewidth = 1.2, figsize=(12,8)) # histograms
```

```
array([[<Axes: title={'center': 'sepal length (cm)'}>,  
       <Axes: title={'center': 'sepal width (cm)'}>],  
       [<Axes: title={'center': 'petal length (cm)'}>,  
       <Axes: title={'center': 'petal width (cm)'}>]], dtype=object)
```

```
plt.show()
```



```
sns.scatterplot(x='sepal length (cm)', y = 'sepal width (cm)', hue = 'species',  
                data = iris_df)  
plt.savefig('iris_scatter.png')  
plt.show()
```



```
scaler = StandardScaler()
iris_scaled = scaler.fit_transform(iris_df.iloc[:, :-1]) # scale data frame
```

```
X_train, X_test, y_train, y_test = train_test_split(iris_scaled, iris.target,
                                                    test_size = 0.3,
                                                    random_state = 42)

model = LogisticRegression()
model.fit(X_train, y_train) # fit the model
```

```
LogisticRegression()
```

```
y_pred = model.predict(X_test) # test the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
```

Accuracy: 1.00