

Lab 6

Conie O'Malley

2025-02-23

Table of contents

0.1	Deliverable 1: Get your working directory and paste below	2
1	Part 1: Building and Using Word Embeddings	2
1.1	Deliverable 1: Load the data and inspect the first few rows	2
1.2	Deliverable 2: Preprocess the data	3
1.3	Deliverable 3: Create a Vocabulary and Term Co-Occurrence Matrix	3
1.4	Deliverable 4: Fit the GloVe Model to the TCM	3
1.5	Deliverable 5: Explore the word embeddings.	4
1.6	Deliverable 6: Find Words Similiar to “king”	4
2	Part 2: Building Simple Bi-Gram Language Models	5
2.1	Deliverable 7: Collect and Prepare Your Data	5
2.2	Deliverable 8: Calculate the Frequency of Bigrams	5
2.3	Deliverable 9: Calculate the Probability of Bigrams	5
2.4	Deliverable 10: Use the Bigram Model to Predict the Next Word	5
3	Part 3: Word Embeddings in Python	6
3.1	Deliverable 11: Load and Prepare the Data in Python	6
3.2	Deliverable 12: Train a Word2Vec model using gensim	7
3.3	Deliverable 13: Explore the word embeddings	7
3.4	Deliverable 14: Find similar words to a “king”	7
3.5	Deliverable 15: Perform Analogies	7

```
# install packages
packages <- c("word2vec", "text2vec", "magrittr")

for (i in packages) {
```

```

if (!requireNamespace(i, quietly = TRUE)) {
  renv::install(i)
}
library(i, character.only = TRUE) # Load the package
}

```

Warning: package 'word2vec' was built under R version 4.3.3

Warning: package 'text2vec' was built under R version 4.3.3

0.1 Deliverable 1: Get your working directory and paste below

```
getwd()
```

```
[1] "/Users/coniecakes/Library/CloudStorage/OneDrive-Personal/001. Documents - Main/023. Proj
```

1 Part 1: Building and Using Word Embeddings

1.1 Deliverable 1: Load the data and inspect the first few rows

```

data("movie_review")
head(movie_review)

```

	id	sentiment
1	5814_8	1
2	2381_9	1
3	7759_3	0
4	3630_4	0
5	9495_8	1
6	8196_8	1

```

1
2
3
4
5
6

```

```
3 The film starts with a manager (Nicholas Bell) giving welcome investors (Robert Carradine)
```

1.2 Deliverable 2: Preprocess the data

```
tokens <- movie_review$review %>%  
  tolower() %>%  
  text2vec::word_tokenizer()
```

1.3 Deliverable 3: Create a Vocabulary and Term Co-Occurrence Matrix

```
it <- text2vec::itoken(tokens, progressbar = FALSE)  
vocab <- text2vec::create_vocabulary(it)  
vectorizer <- text2vec::vocab_vectorizer(vocab)  
  
tcm <- text2vec::create_tcm(it, vectorizer, skip_grams_window = 5L)
```

1.4 Deliverable 4: Fit the GloVe Model to the TCM

```
glove_model <- text2vec::GlobalVectors$new(rank = 50, x_max = 10)  
word_vectors <- glove_model$fit_transform(tcm, n_iter = 20)
```

```
INFO [13:12:21.791] epoch 1, loss 0.1505  
INFO [13:12:23.040] epoch 2, loss 0.0973  
INFO [13:12:24.242] epoch 3, loss 0.0830  
INFO [13:12:25.408] epoch 4, loss 0.0751  
INFO [13:12:26.603] epoch 5, loss 0.0694  
INFO [13:12:27.794] epoch 6, loss 0.0651  
INFO [13:12:28.985] epoch 7, loss 0.0615  
INFO [13:12:30.159] epoch 8, loss 0.0585  
INFO [13:12:31.339] epoch 9, loss 0.0560  
INFO [13:12:32.527] epoch 10, loss 0.0538  
INFO [13:12:33.700] epoch 11, loss 0.0518  
INFO [13:12:34.873] epoch 12, loss 0.0501  
INFO [13:12:36.042] epoch 13, loss 0.0486  
INFO [13:12:37.220] epoch 14, loss 0.0472  
INFO [13:12:38.389] epoch 15, loss 0.0460  
INFO [13:12:39.609] epoch 16, loss 0.0449  
INFO [13:12:40.786] epoch 17, loss 0.0439  
INFO [13:12:41.984] epoch 18, loss 0.0430
```

```
INFO [13:12:43.188] epoch 19, loss 0.0422
INFO [13:12:44.378] epoch 20, loss 0.0414
```

1.5 Deliverable 5: Explore the word embeddings.

```
king_vector <- word_vectors["king", , drop = FALSE]
print(king_vector)
```

```
      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
king 0.2048331 0.1033692 -0.5105069 0.3859732 -0.1183013 0.1526028 0.1755947
      [,8]      [,9]     [,10]     [,11]     [,12]     [,13]     [,14]
king 0.7517204 0.209195 -0.2875485 -0.3756137 -0.1118976 -0.3348667 -0.4776994
      [,15]     [,16]     [,17]     [,18]     [,19]     [,20]     [,21]
king 0.1896476 0.2930924 0.02056652 0.01906195 -0.08077735 0.05754792 -0.195102
      [,22]     [,23]     [,24]     [,25]     [,26]     [,27]     [,28]
king 0.2521377 -0.3297842 0.1747278 -0.4108376 0.4921499 0.2512966 0.3926649
      [,29]     [,30]     [,31]     [,32]     [,33]     [,34]
king -0.008526148 -0.1595909 0.02619935 -0.007153479 -0.3324712 -0.1743605
      [,35]     [,36]     [,37]     [,38]     [,39]     [,40]     [,41]
king -0.0318132 0.3163234 0.1759345 -0.268418 0.424578 0.03814595 -0.3217611
      [,42]     [,43]     [,44]     [,45]     [,46]     [,47]
king -0.6259316 0.0884071 -0.1251529 -0.009555229 0.2984651 -0.5912164
      [,48]     [,49]     [,50]
king 0.6793077 0.07639147 0.6631689
```

1.6 Deliverable 6: Find Words Similiar to “king”

```
cos_sim <- text2vec::sim2(x = word_vectors, y = king_vector, method = "cosine", norm = "l2")
head(sort(cos_sim[,1], decreasing = TRUE), 5)
```

```
      king government      paul      hitler      thin
1.0000000 0.5866684 0.5761881 0.5599212 0.5326028
```

2 Part 2: Building Simple Bi-Gram Language Models

2.1 Deliverable 7: Collect and Prepare Your Data

```
book <- gutenbergr::gutenberg_download(158)
```

Determining mirror for Project Gutenberg from <https://www.gutenberg.org/robot/harvest>

Using mirror <http://aleph.gutenberg.org>

```
bigrams <- book %>%  
  tidytext::unnest_tokens(bigram, text, token = "ngrams", n = 2)
```

2.2 Deliverable 8: Calculate the Frequency of Bigrams

```
bigrams_separated <- bigrams %>%  
  tidyr::separate(bigram, c("word1", "word2"), sep = " ")  
bigram_counts <- bigrams_separated %>%  
  dplyr::count(word1, word2, sort = TRUE)
```

2.3 Deliverable 9: Calculate the Probability of Bigrams

```
word1_counts <- bigrams_separated %>%  
  dplyr::count(word1, sort = TRUE) %>%  
  dplyr::rename(total = n)  
  
bigram_probabilities <- bigram_counts %>%  
  dplyr::left_join(word1_counts, by = "word1") %>%  
  dplyr::mutate(probability = n/total)
```

2.4 Deliverable 10: Use the Bigram Model to Predict the Next Word

```
predict_next_word <- function(current_word) {
  bigram_probabilities %>%
    dplyr::filter(word1 == current_word) %>%
    dplyr::arrange(desc(probability)) %>%
    utils::head(5)
}
```

```
predict_next_word("mr")
```

```
# A tibble: 5 x 5
  word1 word2      n total probability
  <chr> <chr>   <int> <int>      <dbl>
1 mr    knightley 237 1018    0.233
2 mr    elton      171 1018    0.168
3 mr    weston     125 1018    0.123
4 mr    woodhouse   94 1018    0.0923
5 mr    frank       48 1018    0.0472
```

3 Part 3: Word Embeddings in Python

```
#import tensorflow as tf
#import torch
#import keras
import nltk
#from transformers import pipeline
from nltk.corpus import movie_reviews
from gensim.models import Word2Vec
nltk.download("movie_reviews")
```

True

3.1 Deliverable 11: Load and Prepare the Data in Python

```
documents = [list(movie_reviews.words(fileid)) for fileid in movie_reviews.fileids()]
```

3.2 Deliverable 12: Train a Word2Vec model using gensim

```
model = Word2Vec(sentences=documents, vector_size=50, window=5, min_count=2, workers=4)
```

3.3 Deliverable 13: Explore the word embeddings

```
king_vector = model.wv["king"]  
print(king_vector)
```

```
[-0.2856428 -0.17786421  0.2799243  0.15504603  0.04463694  0.43373644  
 1.7166221 -0.5387969 -0.05513414 -0.35213506  0.5866303 -1.0010692  
 1.0382105 -0.2377805  0.96553653 -0.2529933 -0.06705428  0.00373106  
-1.0715698 -1.4174579  1.0173542  0.24079442  1.856158 -1.100083  
 0.79750067 -0.38062474  0.18164986 -0.47823712 -1.079246 -0.16276346  
 0.28239533 -0.8364547  0.03266196  1.0683049 -0.16547683  0.3906276  
-0.31353572  1.2597526 -0.9338115 -0.6319033 -0.36353156 -0.02366525  
 0.13028109  0.29832596  0.69743013  0.55217856  0.01818331  0.09821928  
-0.09170736  0.847732  ]
```

3.4 Deliverable 14: Find similar words to a “king”

```
similar_words = model.wv.most_similar("king", topn=5)  
print(similar_words)
```

```
[('jackson', 0.8497037887573242), ('jane', 0.8455929756164551), ('ryan', 0.8443729877471924)]
```

3.5 Deliverable 15: Perform Analogies

```
result = model.wv.most_similar(positive=["woman", "king"], negative=["man"], topn=1)  
print(result)
```

```
[('jane', 0.8494290113449097)]
```

Transferred Part 4 to a separate ipynb file